


Cassandra

A distributed NoSQL Database

- Open source + Enterprise versions exists
- =Amazon DynamoDB + Google Bigtable
- NoSQL: Lacks full ACID transaction support; Lightweight transactions, f.ex no Joins. Row-atomic updates.
- Configurable replication & availability strategy
- No special master node (multimaster db replication)
- Partitioned key-oriented queries
- Gossip protocol
- Runtime network topology changes & recovery




(<https://cassandra.apache.org/doc/latest/cassandra/architecture/overview.html>)

Cassandra CQL	vs SQL
CREATE KEYSPACE myDatabase WITH replication = {'class': 'SimpleStrategy', 'replication_factor': 1};	CREATE DATABASE myDatabase;
USE myDatabase;	--"---
CREATE TABLE IF NOT EXISTS myTable (id INT PRIMARY KEY); -- (synonyms in cql: COLUMNFAMILY=TABLE) NB table need primary key in CQL.	--"---
ALTER TABLE myTable ADD myField INT;	--"---
CREATE INDEX myIndex ON myTable (myField);	--"---
INSERT INTO myTable (id, myField) VALUES (1, 7);	--"---
SELECT * FROM myTable WHERE myField = 7;	--"---
SELECT COUNT(*) FROM myTable;	--"---
DELETE FROM myTable WHERE myField = 7;	--"---

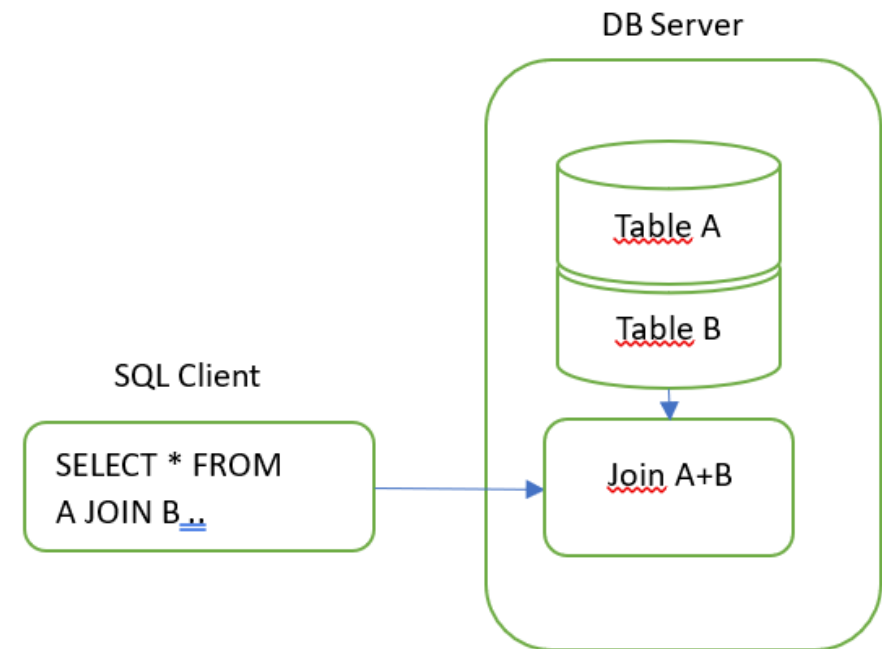
CQL	SQL
<p>- No support for things like JOIN, GROUP BY, or FOREIGN KEY.</p> <p>Leaving these features out is important because it makes writing and retrieving data from Cassandra much more efficient.</p>	<p>JOIN, GROUP BY, FOREIGN KEY</p>
<p>Writes are cheap. Write everything the way you want to read it.</p> <p>CQL does not perform a read while inserting. Without a read, there is no way to know if the data being inserted is replacing an existing record. This means that both inserts and updates are extremely fast.</p>	
<p>UPDATE myTable SET myField = 2 WHERE id = 6;</p> <p>- However, if the row does not exist, it will still get created. Similarly as unintuitive, an INSERT statement will actually replace data if it exists. In where-clause, only primary key column can be used.</p> <p>Under the hood, INSERT and UPDATE are treated the same by Cassandra ("Upserts"), except for Counter columns/tables. Both INSERT and UPDATE require complete PRIMARY KEY.</p>	<p>--"</p>
<p>Transaction Control Language (TCL) - Not in CQL</p>	<p>COMMIT – It saves the work done</p> <p>SAVEPOINT – It identifies a point in a transaction to which you can later roll back</p> <p>ROLLBACK – It restores database to original since the last COMMIT</p>
<p>Data Retrieval/Query Language (DRL/DQL): Simple transactions</p> <p>(Relation between database objects is not possible):</p> <p>- Where clause: only on primary key or secondary indexes!</p> <p>- Can use only AND operator, There are no OR and NOT operators.</p>	<p>Data Retrieval/Query Language (DRL/DQL): Full transactions.</p>

Join?

CQL




SQL




Rollback?

CQL



SQL



CQL Primary Key

```
CREATE TABLE pk_table(  
    p1 <type>, -- partition key  
    p2 .., -- partition key  
    c .., -- cluster key  
    f .., -- cluster key  
    f .., -- non-key column  
    PRIMARY KEY ((p1,p2), c, f)  
);
```

- Key order matters! Left-to-right priority
- Partition keys required.

keyspace_name

table_name


column_name_1	CQL Type	K	←-----	Partition key column
column_name_2	CQL Type	C↑	←-----	Clustering key column (ASC)
column_name_3	CQL Type	C↓	←-----	Clustering key column (DESC)
column_name_4	CQL Type	S	←-----	Static column
column_name_5	CQL Type	IDX	←-----	Secondary index column
column_name_6	CQL Type	++	←-----	Counter column
[column_name_7]	CQL Type		←-----	List collection column
{column_name_8}	CQL Type		←-----	Set collection column
<column_name_9>	CQL Type		←-----	Map collection column
column_name_10	UDT Name		←-----	UDT column
(column_name_11)	CQL Type		←-----	Tuple column
column_name_12	CQL Type		←-----	Regular column


Single partition operation

The **single partition batch operations are unlogged by default** and thus, do not suffer from performance penalties due to logging.

The below diagram depicts the single partition batch request flow from the coordination node *H* to the partition node *B* and its replication nodes *C*, *D*:

<https://www.baeldung.com/java-cql-cassandra-batch>





= hash(PrimKey)
Partition#

Nodes →

	1	2	3	4	5	Replicas:3
1	SE	SE	SE			
2		DE UK	DE UK	DE UK		
3			NO	NO	NO	
4	DK			DK	DK	
5	UK	UK			UK	

Keyspace->

	1	2	3	4	Replicas:1
1	SE				
2		DE UK			
3			NO		
4				DK	

	1	2	Replicas:2
1	SE	SE	
2	DE UK	DE UK	
3	NO	NO	
4	DK	DK	

	1	2	Replicas:1
1	SE		
2		DE UK	
3	NO		
4		DK	

	1	Replicas:1
1	SE	
2	DE UK	
3	NO	
4	DK	

	1	2	3	4	
1	SE, DE UK	SE, DE UK			
2		NO, DK	NO, DK		

userid	firstname	lastname	email
3b4c48a7-13b5-4aba-9f0a-dc75ded08a99	Barney	Rubble	rubble@hotmail.com
2506535a-4999-438d-8682-d5a739596343	Fred	Flinstone	fred@gmail.com
31e24f9d-0d27-4143-82b1-fa1a4268d028	Joe	Rockhead	joer@yahoo.com
57bae20b-3694-4975-a274-db5e856d24ab	Wilma	Flinstone	wilma@bedrock.com

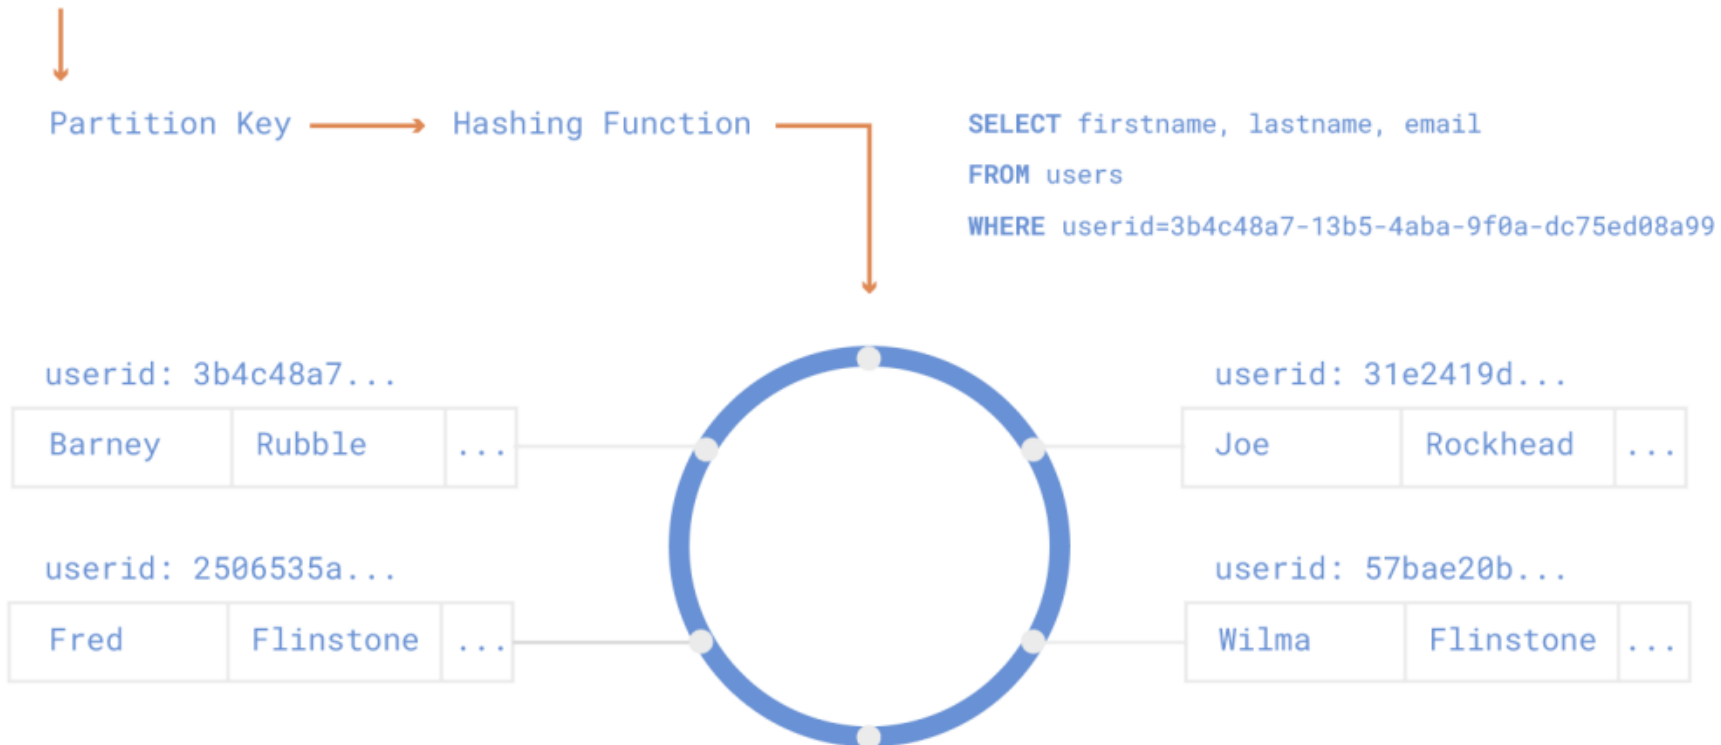


FIGURE 1 How Cassandra stores data